

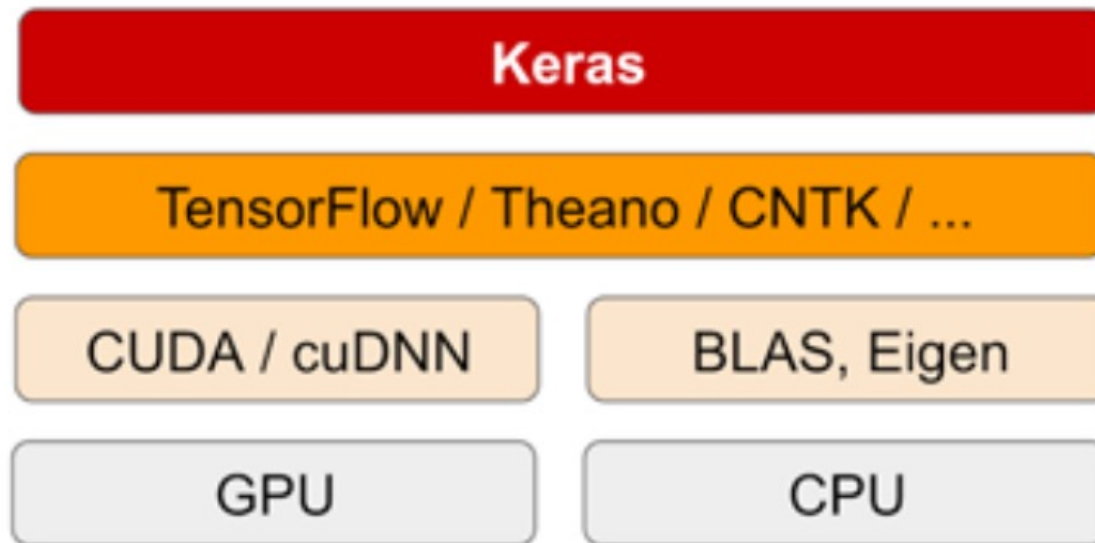


Deep Learning

Fall 2022

Introduction to KERAS

What is Keras?



- Keras is a high-level API providing easy to use elements for deep learning
- Can work with several backends
- Programs can easily deployed on CPUs, GPUs without changing the code

Who makes Keras? Contributors and backers

 **633** contributors



The Keras user experience

- Keras API is easy to understand.
 - Keras follows best practices for reducing cognitive load: it offers consistent & simple APIs, it minimizes the number of user actions required for common use cases, and it provides clear and actionable feedback upon user error.
- Easy to learn.
 - As a Keras user, you are more productive, allowing you to try more ideas than your competition, faster -- which in turn helps you win machine learning competitions.
- This ease of use does not come at the cost of reduced flexibility:
 - Keras integrates with lower-level deep learning languages (in particular TensorFlow), it enables you to implement anything you could have built in the base language. In particular, as `tf.keras`, the Keras API integrates seamlessly with your TensorFlow workflows.

Multi-Backend, Multi-Platform

- Develop in Python, R
 - On Unix, Windows, OSX
- Run the same code with...
 - TensorFlow
 - Cognitive Toolkit (CNTK) - Microsoft
 - Theano
 - MXNet
 - PlaidML
 - ??
- Run on CPU, NVIDIA GPU, AMD GPU, TPU...



How to use Keras: An introduction

Three API Styles

- **The Sequential Model**

- Dead simple
- Only for single-input, single-output, sequential layer stacks
- Good for 70+% of use cases

- **The functional API**

- Like playing with Lego bricks
- Multi-input, multi-output, arbitrary static graph topologies
- Good for 95% of use cases

- **Model subclassing**

- Maximum flexibility
- Larger potential error surface

The Sequential API

```
from tensorflow import keras
from tensorflow.keras import layers

model = keras.Sequential()
model.add(layers.Dense(20, activation='relu', input_shape=(10,)))
model.add(layers.Dense(20, activation='relu'))
model.add(layers.Dense(20, activation='softmax'))

model.fit(x, y, epochs=10, batch_size=32)
```


Options for Layers

- Core Layers
- Convolutional Layers
- Pooling Layers
- Locally-connected Layers
- Recurrent Layers
- Reshape Layers
- Dropout Layers
- Merge Layers
- Normalization Layers
- Noise layers

Options for Layers

- The core layers perform the most basic operations
- They are enough to built FFN networks
- Core Layers
 - Input Layers
 - Dense Layers
 - Activation Layer
 - Embedding Layers
 - Masking layers
 - Lambda Layers

Dense Layer

```
tensorflow.keras.layers.Dense(units, #Number of units in the layer
    activation=None,                 #Standard: use linear output
    use_bias=True,                   #Add a bias vector
    kernel_initializer='glorot_uniform', #How to initialize the weights
    bias_initializer='zeros',         #How the biases
    kernel_regularizer=None,         #For example, apply L2 regularization
    bias_regularizer=None,          #For example, apply L2 regularization
    activity_regularizer=None,       #For example, apply L2 regularization
    kernel_constraint=None,         #For example, non-negative constraint
    bias_constraint=None           #For example, non-negative constraint
)
```

Activation Function

```
model.add(Dense(64))  
model.add(Activation('tanh'))  
#This is equivalent to:  
model.add(Dense(64, activation='tanh'))
```

- Available Activations:
 - softmax
 - elu: (Exponential linear unit.)
 - x if $x > 0$ and $\alpha * (\exp(x) - 1)$ if $x < 0$.
 - selu: Scaled Exponential Linear Unit
 - softplus
 - $\log(\exp(x) + 1)$
 - relu
 - $\text{relu}(x, \alpha=0.0, \text{max_value}=\text{None})$
 - sigmoid
 - tanh

Compiling the Model

```
from tensorflow.keras import optimizers

model.compile(
    optimizer=optimizers.RMSprop(lr=0.001),
    loss='binary_crossentropy',
    metrics=['accuracy']
)
```

- Before training a model, you need to configure the learning process, which is done via the compile method, defining
 - **An optimizer.** This could be the string identifier of an existing optimizer (such as rmsprop or adagrad), or an instance of the Optimizer class
 - **A loss function.** This is the objective that the model will try to minimize. It can be the string identifier of an existing loss function, or it can be an objective function.
 - **A list of metrics.** A metric could be the string identifier of an existing metric or a custom metric function.

Examples: Compiling Models

```
# For a multi-class classification problem
```

```
model.compile(optimizer='adam',  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

```
# For a binary classification problem
```

```
model.compile(optimizer='rmsprop',  
              loss='binary_crossentropy',  
              metrics=['accuracy'])
```

```
# For a mean squared error regression problem
```

```
from tensorflow.keras import optimizers  
model.compile(optimizer=optimizers.RMSprop(lr=0.0025),  
              loss='mse')
```

What does compile do?

- Compile defines the loss function, the optimizer and the metrics. That's all.
 - It has nothing to do with the weights and you can compile a model as many times as you want without causing any problem to pretrained weights.
 - You need a compiled model to train (because training uses the loss function and the optimizer). But it's not necessary to compile a model for predicting.
 - Do you need to use compile more than once? Only if:
 - You want to change one of these:
 - Loss function
 - Optimizer / Learning rate
 - Metrics
 - You loaded (or created) a model that is not compiled yet. Or your load/save method didn't consider the previous compilation.
- Consequences of compiling again:
 - If you compile a model again, you will lose the optimizer states.

Loss Functions

- mean_squared_error
- mean_absolute_error
- mean_absolute_percentage_error
- mean_squared_logarithmic_error
- **binary_crossentropy**
- **categorical_crossentropy**
- sparse_categorical_crossentropy
- ...

Metrics

- Can be any of the loss functions
- Some standard metrics like
 - F1
 - Precision
 - Recall
 - accuracy

Train the Model

```
model.fit(x=None, y=None,
         batch_size=None,
         epochs=1,
         verbose=1,
         callbacks=None,
         validation_split=0.0,
         validation_data=None,
         shuffle=True,
         class_weight=None,
         sample_weight=None,
         ...
)
```

Input and desired outcome
Number of samples per gradient update. If none, it defaults to 32
Number of runs over the complete x and y
Verbosity mode. 0 = silent, 1 = progress bar, 2 = one line per epoch.
List of functions to call during training
Part of dataset set aside for validating
Validation dataset, tuple (x_val, y_val)
shuffle the training data before each epoch
Give some classes more/less weight
Give some samples more/less weight

In Context

```
model.compile(optimizer='rmsprop',  
              loss='binary_crossentropy',  
              metrics=['acc'])  
  
history = model.fit(x=partial_x_train, y=partial_y_train,  
                   epochs=20,  
                   batch_size=512,  
                   validation_data=(x_val, y_val))
```

The History Object

- Note that the call to `model.fit()` returns a History object. This object has a member `history`, which is a dictionary containing data about everything that happened during training.

```
>>> history_dict = history.history
>>> history_dict.keys()
[u'accuracy', u'loss', u'val_accuracy', u'val_loss']
```

- The dictionary contains four entries: one per metric that was being monitored during training and during validation.
- You can now plot these to get Information about your performance

Plotting the training and validation loss

```
import matplotlib.pyplot as plt


history_dict = history.history
loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']

epochs = range(1, len(loss_values) + 1)

# 'bo' is for blue dot, 'b' is for solid blue line
plt.plot(epochs, loss_values, 'bo', label='Training loss')
plt.plot(epochs, val_loss_values, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```

Plotting the training and validation loss



Practical Recommendations

- For lower data amounts, you should train smaller and shallower networks in order to prevent overfitting
- Preprocessing
 - Take small values - Typically, most values should be in the 0–1 range.
 - Be homogenous- That is, all features should take values in roughly the same range.

Load and Save models

- You save a Keras model into a single HDF5 file which will contain:
 - the architecture of the model, allowing to re-create the model
 - the weights of the model
 - the training configuration (loss, optimizer)
 - the state of the optimizer, allowing to resume training exactly where you left off.

```
from tensorflow.keras.models import load_model

model.save('my_model.h5') # creates a HDF5 file 'my_model.h5'
del model                 # deletes the existing model

# returns a compiled model
# identical to the previous one
model = load_model('my_model.h5')
```